From Public ED Data to Southlake Future-State Planning Scenarios

An agentic workflow — with reasoning and analysis technically supported by theFinlyApp — th[illegible] department [illegible] cts, then [illegible] anyone [illegible]

CORE MESSAGE FOR JUDGES

The emphasis is on **how the system thinks and acts**, not just the data it produces. The agentic reasoning and analysis engine is technically supported by **theFinlyApp** — powering every step from intent framing to quality evaluation.

1

Intent Agent

Interprets the stakeholder's planning goal and frames it as a specific operational question for synthesis.

LLM REASONING

→

2

Profile Agent

Analyzes dataset quality, identifies sensitive columns, flags missingness, and recommends focus areas.

LLM REASONING

→

3

Strategy Agent

Decides synthesis approach, target rows, and constraints. This is the key decision step — not scripted.

LLM DECIDES

→

4

Evaluate Agent

Scores fidelity, privacy, and utility. Can trigger a retry if quality is below threshold — self-correction loop.

LLM EVALUATES

→

5

Narrative Agent

Crafts planning summaries tailored to the scenario, stakeholder, and evaluation results.

LLM NARRATES

INPUT

Curated CDC/NCHS NHAMCS 2022 public ED dataset — 13,085 visits, 21 operational columns

SOUTHLAKE LENS

Distributed campus routing · Community diversion · Regional growth · ED surge

OUTPUT

Synthetic dataset + evaluation metrics + caution report + pitch artifacts + downloadable ZIP

WHAT MAKES IT AGENTIC

AI reasons and decides at each step · Strategy adapts to data · Evaluate can trigger retries · Powered by theFinlyApp reasoning engine

WHAT THE STUDIO ACTUALLY DOES

A Planning Studio, Not Just a Data Generator

Saved demo reliability, visible AI reasoning, Southlake-specific scenarios, and exportable planning artifacts in one screen.

Southlake Scenarios

Four planning scenarios aligned to distributed-network strategy: campus routing, community diversion, growth, and surge.

Visible AI Reasoning

Reasoning trace at every agent step — judges can see what the AI decided and why. Powered by theFinlyApp.

Evaluation Dashboard

Fidelity, privacy, utility scores plus source-vs-synthetic preview in one visible screen.

Export & Narrative

Auto-generated methodology, features, cautions + downloadable ZIP for workshops and presentations.

LIVE DEMO OUTPUT — DISTRIBUTED CAMPUS ROUTING

FIDELITY	PRIVACY	UTILITY	EXACT OVERLAP
87.8	100.0	92.7	0%

AGENT REASONING TIMELINE — POWERED BY THEFINLYAPP

INTENT AGENT
"Network operations needs to test routing across campus, observation, and community pathways..."

PROFILE AGENT
"13,085 rows, 21 columns. 42% pain_scale missingness reflects real clinical documentation..."

STRATEGY AGENT
"Selecting GaussianCopula over CTGAN — preserves multivariate flow relationships critical for routing..."

EVALUATE AGENT
"Fidelity 87.8 strong for planning use. Categorical similarity preserved, zero-inflated columns excluded for accuracy..."

NARRATIVE AGENT
"Framing as safe experimentation sandbox — emphasizing innovation velocity with governance boundaries..."

Scenario: Distributed Campus Routing · 13,085 synthetic rows · Stakeholder: Network Operations

DOWNLOAD ZIP

What the Studio Is Safe For, and What It Is Not

Synthetic data is powerful for planning and experimentation. But it is not local operational truth, and it should not be used without governance and validation.

Good For

- Planning workshops and innovation sprints — test routing, capacity, and community-diversion ideas safely
- Scenario-driven discussions about Southlake's distributed-campus future state
- Prototyping dashboards, workflows, and analytics without real patient records
- Pitch-day storytelling backed by visible metrics, AI reasoning, and exportable artifacts

Not Ready For

- Clinical care, staffing commitments, reimbursement, or patient-level action
- Claims about Southlake's real patient mix from a U.S. public-use source
- A full hospital digital twin or validated multi-table operational model
- Production use without governance, privacy review, and validation on governed local data

BEST NEXT STEP

Validate this workflow with governed Southlake data, operational owners, and formal privacy review. The agentic architecture is ready — it needs real local data to become a real planning tool.

WHY THIS SLIDE MATTERS

In healthcare, the team that says "here's what we don't know yet" is more trustworthy than the team that claims everything is solved. This caution boundary is a feature, not a weakness.